Expert-routing proxy

Recorded median latency · lower is faster

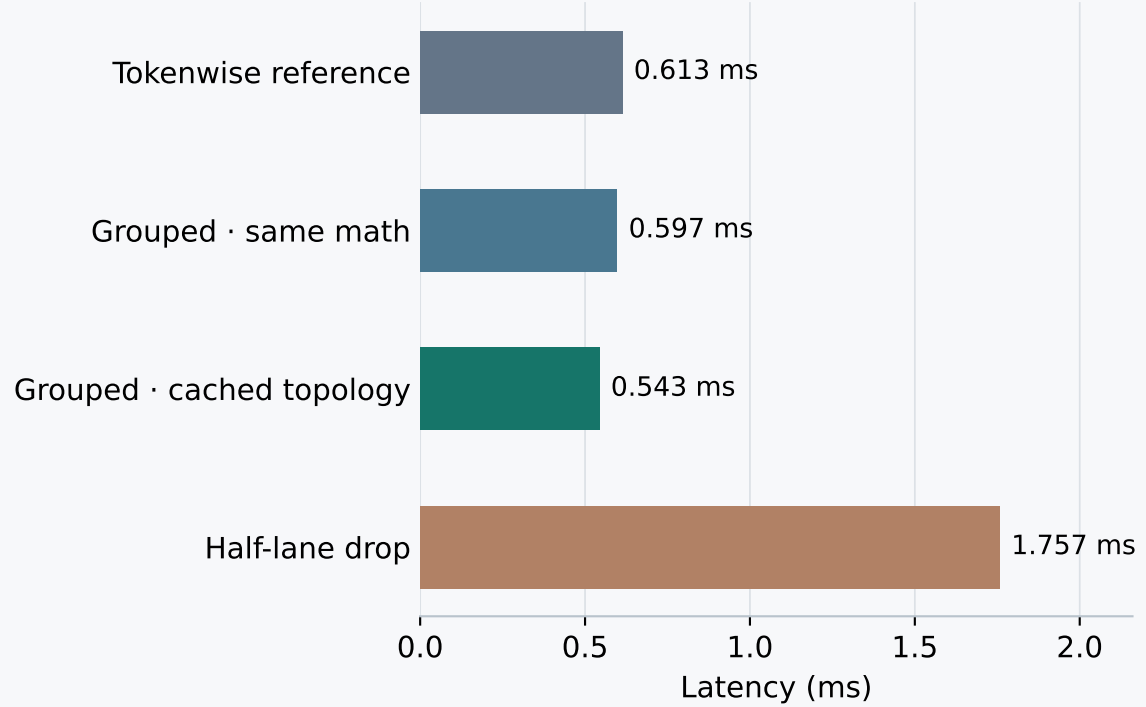

HP15 · GPU buffer reuse

Fresh-buffer / pooled latency · baseline run

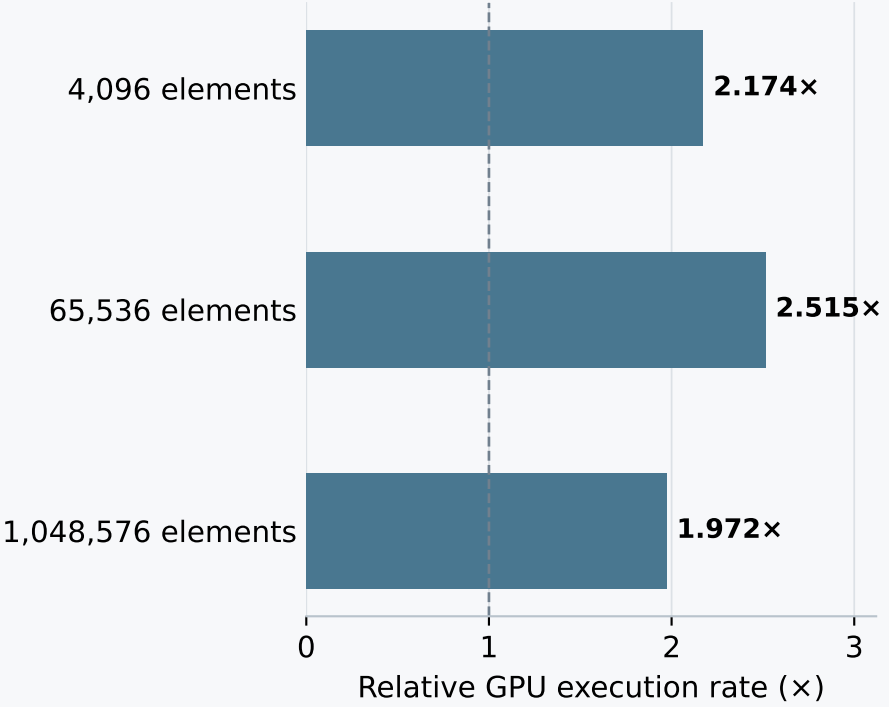

Separate workloads and baselines. Source: data/expert-routing.csv and data/hp15-baseline.json.
HP15 ratios compare two GPU allocation paths; CPU comparisons and output fidelity are provided in RESULTS.md.